

Volumetric Optics Simulator Design

https://gemini.google.com/app/5a76969bf8580428

User prompt: NO, first of all tell me, what's the best effort:product ratio here? web native app low level app (native app but with low level language like rust or C++ etc)

Response: The clear winner for effort-to-product ratio is **Web with WebGL / WebGPU (via a framework like PixiJS, Three.js, or raw Shaders)**.

Platform	Effort (Dev Time & Overhead)	Performance Ceiling	Distribution & Iteration	Verdict
Web (WebGL/WebGPU)	Low–Medium Instant setup, no build-tool hell, rich 2D UI libraries for sliders/menus.	Very High Runs directly on the GPU. A single shader handles millions of pixels effortlessly on integrated graphics.	Instant One URL to share; runs cross-platform without installs or security warnings.	 Best Ratio
Native (Engine: Godot/Unity)	Medium Ready-made 2D physics, UI, and viewport pipelines out of the box.	High Low overhead, good memory management.	Medium Need to compile binaries or export to web (WASM exports can be bulky).	Solid Alternative
Low-Level Native (C++/Rust + Vulkan/OpenGL)	Extremely High You build windowing, UI, math libraries, pipeline management, and memory safety yourself.	Maximum Near zero CPU overhead, maximum control over thread dispatching.	Painful OS-specific packaging, driver quirks, manual installation required.	 Worst Ratio

Why Web Wins for this Specific Project

- The Bottleneck Isn't the CPU:** Optics and continuous light haze are fundamentally **fill-rate and GPU-bound problems**, not CPU-bound. Whether you write a fragment shader in C++, Rust, or WebGL, the GPU executes the exact same compiled SPIR-V/GLSL shader bytecode. C++ won't make a GPU fragment shader run any faster than WebGL will.
- Rapid Prototyping of UI & Controls:** Half of an interactive optics simulator is the control interface (dragging lenses, rotating slits, tweaking refraction indices). Building draggable knobs, sliders, and color pickers in HTML/CSS takes minutes; doing the same in C++ with ImGui or custom UI takes days.
- Budget Hardware Compatibility:** WebGL 2.0 has essentially 100% support across modern browsers and runs smoothly on low-end laptops (Intel HD/UHD graphics) as long as you use **screen-space fragment shaders or 2D polygon meshes** instead of raw CPU loops.

The Recommended Stack

- Rendering:** Raw WebGL2 fragment shader on a full-screen canvas (or PixiJS if you want easy sprite/element drag-and-drop).
- UI/State:** Vanilla JS or Svelte/React for lightweight parameter sliders.
- Technique:** Screen-space 2D raymarching or analytical 2D mesh projection—keeping all math on the GPU.

User prompt: hmmm, lets start ambitious with a Truly Performance and Visual Ratio balanced, keeping in mind the diverse support for any device so, the only worry here i think is the Worry about continuous Render loops right? because light is light, we can't force it to be clamped without making the whole thing unrealistic... any idea's?

Response: The key to solving the continuous loop problem without sacrificing realism or melting budget hardware is to **stop treating light as an infinite continuous cascade every single frame** and instead take advantage of how real-time rendering cheats wave physics.

Here are the three structural strategies that give you continuous, non-pixelated light beams across low-end mobile and budget laptops alike:

1. Demand-Driven "Dirty" Rendering (Event-Driven vs. 60 FPS Polling)

Light travels at 3×10^8 m/s, which means in a static 2D scene, the light field reaches steady-state **instantly**.

- The Trap:** Running a heavy continuous computation loop 60 times a second even when nothing is moving.
- The Fix:** Only re-render or re-trace light paths when:

1. The user drags/rotates a lens, emitter, or prism.
2. An animated dynamic object (like a pulsating black hole or floating dust) updates.

- When the scene is static, **stop tracing entirely** and hold the frame buffer. CPU and GPU usage drop to near 0%, keeping laptops completely cool.

2. Temporal Accumulation (Progressive Refinement)

Instead of forcing your GPU to render 5,000 rays or 200 raymarch steps in a single 16ms frame:


- **While Moving (Interaction State):** Trace a coarse beam (e.g., 64 ray steps or 100 sample lines) with an ultra-light bilateral blur. This guarantees a locked 60 FPS while dragging.
- **When Resting (Settled State):** Accumulate samples over 8–16 frames into a floating-point render target ($1/2$ or $1/4$ resolution buffer). Within ~200 milliseconds of stopping mouse movement, the beam smooths out into a continuous, crystal-clear volumetric glow.

3. Screen-Space 2D Raymarching with Early Exit (The Math Cheat)

To avoid infinite recursion loops when light bounces inside prisms or loops around a black hole event horizon:

- **Signed Distance Fields (SDFs):** Represent all shapes (slits, prisms, curved glass, circles) as 2D distance functions inside a single fragment shader.
- **Adaptive Step-Size (Sphere Tracing in 2D):** The ray takes large leaps through empty air and only takes microscopic steps near glass boundaries or the high-gravity gradient of a black hole.
- **Energy Threshold Clamping (Not Distance Clamping):** Never clamp by an arbitrary short range. Instead, track the ray's radiance vector (I_R, I_G, I_B). As light reflects and transmits, multiply by Fresnel coefficients ($T \approx 0.96, R \approx 0.04$). As soon as $|I| < 0.005$ (below human eye visibility on an 8-bit monitor), terminate the ray immediately.

Suggested Pipeline Overview



User prompt: don' you think instead of rushing though everything we should talk it out and plan everything first? (Recognising edge cases, Critiquing my own framework by questioning if it will hold for specific scenarios or not, catching bottle necks before execution will save us later regrets) shall we start then?

Response: Spot on. Diving straight into code before pressure-testing the core mechanics is how physics engines end up rewritten three times.

Simulating light as a continuous, volumetric medium with refraction, dispersion, and extreme curvature brings massive theoretical and technical pitfalls.

Here is the breakdown of the critical friction points and edge cases we must resolve first:

1. The Rendering Model (Forward vs. Backward Tracing)

- **The Dilemma:**
 - **Forward (Source → Scene):** The emitter shoots light out. It accurately produces caustics, focal hotspots, and dispersed rainbow beams, but it is expensive to turn millions of points into a continuous screen-space fog without severe noise or overdraw.
 - **Backward (Screen Pixel → Scene):** Every pixel casts a ray back to check if it hits a light source. This yields noise-free full-screen volumetric scattering easily, but it is notoriously terrible at finding tiny, refracted caustics (e.g., light focused through a small convex lens onto a focal point).
- **The Question:** Do we use forward photon-mapping/line-bundling onto an off-screen accumulation texture, or reverse screen-space raymarching?

2. Physical Edge Cases & Branching Explosions

- **Total Internal Reflection (TIR):** Inside a high-index prism ($n \approx 1.5$), rays hitting the boundary beyond the critical angle ($\theta_c = \arcsin(1/n) \approx 41.8^\circ$) don't exit; they reflect infinitely inside the geometry until an angle changes.
- **Ray Branching vs. Budget Limits:** At every single glass interface, a real light beam splits into **Reflected (Fresnel)** and **Refracted (Snell)** components.
 - 1 bounce $\rightarrow$ 2 beams.
 - 4 bounces $\rightarrow$ 16 beams.
 - 10 bounces $\rightarrow$ 1,024 active beams per sample.
- **The Question:** How strictly do we truncate or cull secondary reflection branches without losing realism?

3. Chromatic Dispersion

- White light is a continuous spectrum ($\approx 380\text{ nm}$ to 750 nm).
- If we only sample discrete RGB bands (Red, Green, Blue), a prism will output **three distinct colored lines** instead of a smooth, blended rainbow smeared across space.
- **The Question:** How do we reconstruct a continuous smooth spectral smear efficiently without tracing 50 separate wavelength passes per beam?

4. Non-Euclidean Optics (The Black Hole)

- Normal optics use straight-line segment intersections ($y = mx + b$).
- Gravitational fields bend spacetime continuously, meaning rays follow curved geodesic paths ($d^2x^\mu/d\lambda^2 + \Gamma^\mu_{\alpha\beta} \cdot \dots$).
- Straight-line intersection math completely breaks here; every curved ray requires numerical integration (Euler or Runge-Kutta step-by-step march).
- **The Question:** Can we hybridize straight-line analytic optics for glass/prisms with localized field-marching only within the gravitational radius of the black hole?

Where do you want to begin deconstructing first: **the core light transport paradigm (forward vs. backward)** or **the mathematical limits of ray-branching and dispersion**?

User prompt: lets start with 1 first, tell me all the pros and cons of each

Response: 1. Forward Light Transport (Source $\rightarrow$ Scene)

You simulate photons starting from the emitter slit, tracing them forward as they hit glass surfaces, bounce, split, and travel outward.

Pros	Cons
Effortless Caustics & Focal Points: When a convex lens focuses a beam to a tiny point, thousands of rays naturally converge to the exact same screen coordinates. Brightness ramps up automatically via additive blending.	Screen Reconstruction Noise: Light between rays is completely empty. To make a beam look like a smooth, continuous volume rather than visible discrete needles, you must trace tens of thousands of rays or apply heavy spatial blurs.
Intuitive Chromatic Dispersion: Splitting a ray into multiple wavelengths at a prism boundary maps directly to physical intuition ($n_{\text{blue}} > n_{\text{red}}$).	Severe Overdraw on GPUs: If 10,000 rays overlap to draw a single bright beam, the GPU re-writes the exact same pixels thousands of times, crashing fill rate on budget integrated GPUs.
Natural Energy Conservation: Energy splits cleanly at interfaces using Fresnel reflection/transmission coefficients ($R + T = 1$).	Wasted Work: Any photon that shoots off-screen or gets trapped inside geometry still consumes compute cycles despite contributing nothing to the visible frame.

2. Backward Light Transport (Screen Pixel $\rightarrow$ Scene)

Every screen pixel acts as a sensor, marching backward through the 2D scene along a line of sight or checking if it sits inside a lit volume to evaluate scattered light.

Pros	Cons
Deterministic GPU Load (Fixed Cost): Computation scales strictly with screen resolution (e.g., 1920×1080), never exploding when ray counts or bounce depths surge.	The "Delta Light" Problem (Missing Caustics): A pixel cannot easily "find" a tiny light source that has been refracted through an arbitrary curved lens. Pinpointing where a refracted ray came from backwards requires solving complex inverse optical equations.

Pros	Cons
Naturally Continuous & Noise-Free: Volumetric fog and smooth gradients render natively without visible ray striping, because every neighboring pixel evaluates a continuous field.	Complex Dispersion Math: Splitting colors backward is mathematically inverted; you must query which exact wavelength could have reached the current pixel from the source through all intervening refractions.
Zero Off-Screen Waste: Pixels that are dark or outside the viewport never execute deep intersection loops.	Curved Spacetime Headaches: Tracing a path backward from a pixel into a black hole's warped field to see if it hits a tiny emitter slit is numerically unstable (like trying to thread a needle backwards).

The Comparison

Metric	Forward (Ray Bundling / Photon Map)	Backward (Screen-Space Raymarch / SDF)
Focal Point Realism	Exceptional (Natural convergence)	Poor (Requires hard inverse solving)
Atmospheric "Fuzz" Quality	Requires post-blur / high density	Flawless and smooth natively
Performance Scaling	Scales with Ray Count × Bounces	Scales with Screen Resolution × March Steps
Black Hole Deflection	Trivial (Push ray forward along curve)	Hard (Finding if curved ray hits slit)
Suitability for Budget GPUs	High with 2D additive instancing	High with low-step SDF sphere tracing

The Verdict for Our Goal

Given that the core requirements are **slits, lenses, prism rainbows, and black holes**, pure Backward tracing fails at the physics (finding refracted focal points and black hole lensing from a narrow slit is mathematically brutal).

A Hybrid Forward Method is typically the sweet spot:

- 1. **Forward Pass:** Emit a structured ray bundle forward from the slit to compute the exact optical paths, splits, and caustics into a low-resolution float buffer.
- 2. **Backward / Post Pass:** Apply an energy-preserving bilateral blur / scattering pass over that buffer so the beams blend into a silky, continuous volumetric haze without needing millions of rays.

Does this forward-first hybrid approach make sense for the engine, or do you want to probe the failure modes of backward tracing further?


User prompt: why can't we pick The best of both worlds?

Response: We can—and in rendering, this is known as **Bidirectional Light Transport / Photon Mapping**.

However, combining both is not a magic free lunch. When you bridge forward and backward passes, you inherit the mathematical complexity of connecting the two disjoint paths.

Here is what a true "Best of Both Worlds" architecture looks like, along with the exact technical bottlenecks we have to solve before committing to it.

The Architecture: 2D Volumetric Photon/Beam Mapping



Where This Hybrid Model Wins

- 1. **Perfect Caustics & Focusing:** Convex lenses, prisms, and mirrors use the forward pass, so focal points and chromatic dispersion are physically accurate.
- 2. **Infinite Continuity (No "Ray Lines"):** Instead of trying to brute-force 50,000 rays to fill the visual gaps, the screen pass treats every forward ray as a **continuous 2D segment with a soft radial falloff kernel**. Even 300 rays can produce a seamless, continuous volumetric fog.

3. **Cheap on Budget GPUs:** The GPU doesn't suffer massive overdraw from tens of thousands of transparent lines. It renders the scene in two ultra-lightweight stages.

The Friction Points & Blind Spots of a Hybrid Approach

Before building this, we have to stress-test three distinct failure modes:

1. The Kernel Overlap Problem (Over-Brightness at Intersections)

- When you blur or expand ray segments into volumetric beams, what happens when 10 rays cross each other?
- If you do simple additive blending of blurred segments, the intersection can blow out into an unnatural white blob that is wider than the actual physical beam.
- **The Challenge:** The screen-space filter must normalize energy density so that a focused beam gets sharper and brighter, rather than just fatter and blown out.

2. Black Hole Geodesic Segmentation

- Prisms produce straight line segments, which are trivial to store in a buffer $[x1, y1, x2, y2]$.
- A black hole bends rays into smooth, non-linear curves.
- **The Challenge:** To store a curved geodesic in a segment buffer, the forward pass must discretize the curve into small linear piecewise steps. If the step size is too large, the black hole beam looks like an angular polygon; if too small, memory bandwidth explodes.

3. Wavelength Blending (The Rainbow Reconstruction)

- If Pass 1 splits a ray into 3 discrete wavelengths (Red, Green, Blue), Pass 2 needs to interpolate between them in screen-space so the output is a continuous spectrum, not three distinct stripes with black gaps between them.

The Big Architectural Decision

To make the hybrid method work smoothly on a budget laptop, how do you want Pass 1 to talk to Pass 2?

- **Option A: Texture-Based (Rasterize forward rays to an off-screen low-res HDR buffer, then run a directional/bilateral blur shader over it).**
 - *Pros:* Insanely fast on GPUs, zero complex data structures in JS.
 - *Cons:* Blurs across object edges (light might visually "bleed" through a solid opaque obstacle if the blur kernel is too wide).
- **Option B: Geometry/Segment-Based (Forward pass generates a list of 2D line segments; a custom fragment shader calculates the exact shortest distance from each pixel to all active segments).**
 - *Pros:* Razor-sharp optical boundaries, light never bleeds through walls.
 - *Cons:* Pixel shader complexity scales linearly with the number of segments ($O(\text{Pixels} \times \text{Segments})$), which can lag on low-end integrated graphics if bounce depth is high.

Which of these two approaches do you think holds up better against your constraints?

User prompt: ... Suggestion?

Response: Go with **Option A (Texture-Based: Forward HDR Buffer + Depth-Aware Bilateral Scatter)**, but with a slight tweak to fix the edge-bleeding flaw.

Here is why Option B will fail on budget machines and why modified Option A is the right choice:

Why Option B is a Trap


Option B ($O(\text{Pixels} \times \text{Segments})$) is mathematically pure, but integrated Intel graphics (like Intel UHD / Iris Xe) choke on massive per-pixel loop iterations. If you have 300 rays that split and bounce 3 times, that is $\sim 1,200$ segments. At 1920×1080 , your GPU has to evaluate **2.4 billion distance tests every single frame**. That will bring a budget laptop to 10–15 FPS instantly.

Why Option A Wins (With the "Depth-Mask" Fix)

1. **Fixed Performance Cost:** The forward rasterization is tiny ($O(\text{Segments})$ lines drawn with `gl.LINES`), and the blur pass is a fixed $O(\text{Screen Resolution})$ cost regardless of how many times the light bounces.
2. **Eliminating the Edge-Bleed Flaw:** To prevent light from blurring through opaque prisms, mirrors, or barriers, use a **Depth/Obstacle Mask**:
 - Render the physical glass/mirrors/apertures into a 1-bit stencil or 8-bit mask buffer.
 - When running the bilateral blur, sample the mask: if a blur tap crosses from air into a solid obstacle boundary, **zero out that sample's weight**.

- Light smoothly diffuses through the air, wraps naturally around sharp edges, but terminates cleanly at solid surfaces without any bleeding.

The Recommended 3-Step Pipeline



Next Critical Question: Chromatic Dispersion Math

Now that we have locked down **Option A (Forward HDR Buffer + Masked Bilateral Blur)** as the rendering foundation, we should tackle **Problem 3: The Rainbow Spectrum**.

If we only emit Red, Green, and Blue rays, a prism produces **3 discrete stripes** instead of a continuous rainbow smear.

How should we handle the spectrum?

1. **Spectral Sampling (Emit 7–12 discrete wavelength bins like 400nm, 450nm, 500nm... and convert $\lambda \rightarrow \text{RGB}$ via CIE color matching functions)** — physically accurate, smooth rainbow, slight increase in forward ray count.
2. **RGB Cauchy Interpolation (Emit only RGB, but connect them as triangular colored mesh strips instead of lines)** — extremely cheap, but mathematically trickier when rays cross each other.

Which direction do you want to explore?

User prompt: i think both ways are flawed, this is not light anymore, its just Discrete Coloured Light added together, i dont want that, i want True Divergence and convergence

Response: You are completely right to reject that.


Sampling 3 RGB lines or even 12 colored stripes is a hack. In real life, light is not a collection of discrete colored threads; it is a **continuous radiant energy field over a continuous interval of wavelengths** ($\lambda \in [380 \text{ nm}, 780 \text{ nm}]$).

When a continuous beam hits a prism, it doesn't break into separate lines—the wavefront shears into an **infinite, unbroken angular fan**.

If we want **true physical divergence, continuous smooth gradients, and convergence without any striped ray artifacts**, we have to abandon "ray tracing" and move to **Beam / Wavefront Polygon Transport**.

The Fundamental Shift: Ray Tracing $\rightarrow$ Polyhedral Beam Tracing

Instead of tracing individual 1D lines, we trace **2D Continuous Beam Quadrilaterals (Light Frustums)**.



A beam is defined by its boundary curves/segments. Inside that polygon, the radiant energy is mathematically continuous everywhere.

How "True" Continuous Dispersion Works with Beam Geometry

When a continuous white light beam hits a dispersive interface (prism glass with index $n(\lambda) = A + B/\lambda^2$):

1. **Calculate the Extreme Boundaries:**
 - Compute the refraction vector for the shortest visible wavelength ($\lambda_{\text{violet}} \approx 380 \text{ nm}$, maximum bending angle θ_{max}).
 - Compute the refraction vector for the longest visible wavelength ($\lambda_{\text{red}} \approx 780 \text{ nm}$, minimum bending angle θ_{min}).
2. **Construct the Spectral Polygon Mesh:**
 - The refracted beam is no longer parallel; it fans out into a quadrilateral/triangle bounded between θ_{red} and θ_{violet} .
3. **Continuous Analytic Color Evaluation in the Fragment Shader:**
 - Instead of drawing colored lines, you draw the **dispersed polygon mesh** directly to the screen.

- Inside the fragment shader, for every pixel within that polygon, you calculate its normalized position $u \in [0, 1]$ across the angular fan.
- Map u directly to the continuous wavelength:

$$\lambda(u) = \lambda_{\text{red}} + u \cdot (\lambda_{\text{violet}} - \lambda_{\text{red}})$$

- Convert $\lambda(u)$ directly to CIE XYZ $\rightarrow$ sRGB using the continuous **CIE 1931 color matching functions**.

The result: You render a single geometric polygon, but the GPU evaluates a mathematically infinite, continuous, buttery-smooth rainbow across every single subpixel. Zero striping, zero gaps, zero ray-count overhead.

How Convergence and Caustics Behave

- **Collimated Beam:** Parallel boundaries $\implies$ uniform brightness everywhere inside the quad.
- **Focusing through a Convex Lens:** The top boundary A and bottom boundary B converge toward a focal point (x_f, y_f) .
 - The local width of the beam $w(t) \rightarrow 0$.
 - By energy conservation, the fragment shader computes local intensity as $I(t) = I_0 \cdot \frac{w_0}{w(t)}$.
 - As the beam narrows toward the focal point, the brightness naturally ramps up to blazing white caustics without needing thousands of overlapping lines.
- **Inversion Past Focal Point:** The polygon flips over, crossing itself, and expands out the other side naturally.

Edge Cases of Continuous Beam Tracing

Before we lock this in, let's critique this method for the rest of your mechanics:

1. **Splitting on Corner Vertices:** If a continuous beam hits the sharp corner of a prism (where half the beam enters the glass and half misses and stays in the air), the beam polygon must be clipped/split into two sub-polygons. (A 2D polygon clipping algorithm like Sutherland-Hodgman is required).
2. **The Black Hole Problem:** In a flat Euclidean scene, beam boundaries are straight lines. In a gravitational field, the beam boundaries become curved spline paths.

Does this **Continuous Beam Mesh + Analytic Shader Dispersion** match the standard of realism and physical fidelity you're aiming for?

User prompt: suggestions?

Response: Here is the direct recommendation to solve those two edge cases cleanly without ballooning complexity:

1. The Corner Splitting Issue: Adaptive Boundary Sub-division

Instead of writing a bulky computational geometry clipper (like Sutherland-Hodgman) that can fail on floating-point corner edges, use **Adaptive Wavefront Splitting**:

- **How it works:** A beam isn't just a 2-line box; treat it as an array of boundary points along the wavefront (e.g., initialized with 4–8 evenly spaced anchor points across the aperture slit).
- **When an obstacle corner cuts into the beam:**
 - Detect if adjacent anchor paths hit different surfaces (e.g., Path k hits the prism face, Path $k + 1$ misses and hits the back wall).
 - Subdivide the interval at the boundary to snap the exact split point on the prism corner.
 - Split the beam into **two distinct polygon meshes** at that vertex.
- **Why this works:** It maintains clean, sharp, non-overlapping geometric fans with zero messy polygon-clipping edge bugs.

2. The Black Hole Issue: Curved Triangle Strips (Ribbons)

A black hole deforms a straight beam into a sheared, swirling funnel. Trying to represent that as a single rigid quad breaks down immediately.

- **The Fix (Piecewise Wavefront Ribbons):**
 - Step the anchor points forward in time (Δt) along their geodesic curves using a fast 2nd-order Runge-Kutta (RK2) integrator:

$$\frac{d^2 r}{dt^2} = -\frac{GM}{r^3} r$$

- Connect the time steps into a continuous **triangle strip (ribbon mesh)**.

- As the beam gets closer to the photon sphere ($r = 1.5 r_s$), add a step density factor so the ribbon curves smoothly around the event horizon.
- **The Look:** The GPU renders standard connected triangles, but because the vertices follow curved geodesics, the light field warps, stretches, and wraps around the singularity seamlessly.

3. Summary of the Unified Architecture

Physical Phenomenon	How the Engine Represents It	GPU Render Primitive
Straight Beams (Air / Mirrors)	Continuous 2D Linear Polygon	Single Quad / Convex Polygon with soft edge falloff.
Prism Rainbows	Continuous Angular Fan bounded by $[\lambda_{380}, \lambda_{780}]$	Single Triangular Fan with analytic CIE 1931 color shader.
Lenses (Focusing / Caustics)	Converging/Diverging Polygon	Quad with inverse-width brightness scaling ($I \propto 1/w$).
Black Holes	Geodesic-stepped wavefront paths	2D Curved Triangle Strip (Ribbon).

Where to Focus Next

Now that the core mathematical representation (Wavefront Ribbons & Polygons) is locked in, the next design decision is:

1. **The Core Data Structure:** How a single Beam object represents its wavefront points, energy, and child branches in JavaScript/TypeScript.
2. **The Analytic Color Shader:** The exact GLSL function that takes a continuous normalized float $u \in [0, 1]$ and converts it directly into a physically correct continuous sRGB wavelength gradient.

Which one do you want to break down first?

User prompt: so, Lets In Detail Revise Every Decision we went through? So that we dont get confused about this phase later on?

Response: 1. The Core Objective & Hard Constraints

- **The Visual Goal:** Simulate light as a true continuous volumetric field with natural caustics, smooth dispersion, and fuzzy atmospheric scatter—completely eliminating visible discrete ray lines or colored stripes.
- **The Hardware Constraint:** Must run at a locked 60 FPS on budget laptops and integrated graphics (e.g., Intel UHD/Iris Xe) without freezing the CPU or blowing out the GPU fill-rate.

2. Platform & Engine Architecture

Decision Area	Chosen Path	Why We Rejected the Alternative
Tech Stack	Web (WebGL2 / GLSL + Vanilla JS UI)	Low-level native (C++/Rust + Vulkan) yields zero extra performance for GPU-bound shaders while adding massive build/UI development overhead.
Render Strategy	Event-Driven "Dirty" Rendering	Continuously running full optical passes at 60 FPS on a static scene wastes battery and causes thermal throttling. Frames only re-calculate on interaction.
Atmospheric Haze	Depth-Masked Bilateral Scatter Pass	Standard Gaussian blur bleeds light through solid walls and prisms; a depth/obstacle mask cleanly constrains the scatter to the air volume.

3. Light Transport Model: Beam / Wavefront Polygon Tracing

We abandoned both **discrete forward ray casting** (which creates visible needles and GPU overdraw) and **pure backward screen raymarching** (which fails to find refracted caustics).

Instead, light is represented as **Continuous 2D Polygonal Beams (Wavefront Frustums)**:

- **The Core Primitive:** A beam is a 2D geometric surface bounded by its outer wavefront trajectories.
- **Energy Continuity:** Energy is mathematically continuous throughout the area of the polygon, not concentrated onto discrete 1D line segments.

4. Physics Solutions by Optical Phenomenon

- **Prism Dispersion (True Continuous Rainbow):**

- *Method*: Prisms calculate extreme refraction vectors for boundary wavelengths ($\lambda_{\text{violet}} = 380 \text{ nm}$ and $\lambda_{\text{red}} = 780 \text{ nm}$) using Cauchy's equation ($n(\lambda) = A + \frac{B}{\lambda^2}$).
- *Rendering*: An angular triangular fan is rendered with an analytic fragment shader that continuously maps pixel coordinate $u \in [0, 1]$ to continuous wavelength $\lambda(u)$, converted directly to sRGB via CIE 1931 color-matching equations.
- **Lenses & Caustics (Convergence / Divergence)**:
 - *Method*: Boundary points bend via Snell's Law ($n_1 \sin \theta_1 = n_2 \sin \theta_2$).
 - *Rendering*: As the polygon narrows toward a focal point of width $w(t)$, intensity scales dynamically by energy conservation ($I \propto 1/w$), naturally generating blinding focal caustics.
- **Sharp Corner Splitting**:
 - *Method*: Adaptive Wavefront Subdivision detects when adjacent anchor points strike different object faces, splitting the parent beam cleanly into two independent child polygon meshes.
- **Black Hole Gravitational Lensing**:
 - *Method*: Wavefront anchor points are integrated along curved geodesics using Runge-Kutta 2 (RK2) step-wise field marching.
 - *Rendering*: Stepped points are connected into a **curved 2D triangle strip (ribbon mesh)** that smoothly deforms and orbits the singularity.

5. Execution Roadmap

[Phase 1: Math Foundation] → CIE 1931 Spectral GLSL Shader & Snell/Cauchy Refraction Math
 [Phase 2: Geometry Engine] → Wavefront Beam Polygon generation & Corner Subdivision Logic
 [Phase 3: Curvature Engine] → RK2 Geodesic Ribbon Integrator for Gravitational Fields
 [Phase 4: Post-Processing] → HDR Framebuffer + Depth-Masked Atmospheric Scatter Filter

With the entire system architecture, edge cases, and algorithmic models clearly defined, which phase should we build first?

User prompt: ok, i will begin questioning, decision making and planning session of Phase 2 in other chat, finalised phase 1 then? tell me what we have in phase 2

Response: **Phase 2: The 2D Wavefront Geometry Engine**

Phase 2 focuses on turning optical mathematics into physical 2D polygon meshes and managing how light geometry splits, clips, and branches across scene objects.

Core Components in Phase 2

- **1. Wavefront Data Structures & Scene Graph**
 - Defining the Beam structure: anchor points along the wavefront, energy vector (I_R, I_G, I_B) , bounce generation/depth counter, and parent-child branch trees.
 - Representation of scene obstacles (flat mirror segments, prism boundaries, and curved circular/parabolic lens arcs).
- **2. Analytic Surface Intersections (2D)**
 - **Line-to-Segment**: Ray boundary intersection with flat optical faces, apertures, and walls.
 - **Line-to-Arc / Circle**: Ray boundary intersection with convex/concave lens surfaces, including analytic surface normal calculation for Snell's Law.
- **3. Adaptive Corner Subdivision (The Splitting Algorithm)**
 - Detecting geometric discontinuities when adjacent wavefront anchor points hit different surfaces or miss entirely.
 - Subdividing intervals to isolate the exact geometric corner coordinate of a prism/lens.
 - Splitting the parent beam into independent child polygons without creating visual seams or vertex overlaps.
- **4. Dynamic Vertex Buffer Construction**
 - Packing generated beam polygons into GPU-ready `Float32Array` buffers.
 - Assigning per-vertex attributes: Position (x, y), Normalized Dispersion Coordinate $u \in [0, 1]$, Beam Width / Intensity scale factor, and Energy attenuation.
- **5. Branch Pruning & Total Internal Reflection (TIR)**
 - Trapping and reflecting beams that exceed the critical angle ($\theta > \theta_c$).

- Managing the recursion tree for Fresnel splits (reflected beam vs. transmitted beam) and terminating branches whose radiant intensity drops below the visibility threshold ($I < 0.005$).

Key Questions to Address in Phase 2 Planning

1. **Polygon Topology:** Do we triangulate beams into quad strips on the CPU, or send boundary control points to the GPU and tessellate there?
 2. **Subdivision Budget:** What is the maximum subdivision depth for corner splits to prevent worst-case CPU spikes during rapid mouse dragging?
 3. **Branch Tree Limit:** How many simultaneous active beam branches do we allow before hard-capping the scene tree?
-